

# Introduction to Computer Organization

**DIS 1A – Week 3**

Slides modified from UT WANG

# Agenda

- All about functions !
- Arrays
- Struct and Union
- Midterm tips

# Exercise: Fun with arithmetic

(a) C code

```
1  int arith(int x,  
2          int y,  
3          int z)  
4  {  
5      int t1 = x+y;  
6      int t2 = z*48;  
7      int t3 = t1 & 0xFFFF;  
8      int t4 = t2 * t3;  
9      return t4;  
10 }
```



**Figure 3.8** C and assembly code for arithmetic routine body. The stack set-up and completion portions have been omitted.

Convert this function to x86. Assume that: x at %ebp+8, y at %ebp+12, z at %ebp+16. Recall: `addl src dst`, `imull src dst`, and `andl src dst`.

# Exercise: Fun with arithmetic

(a) C code

```
1  int arith(int x,  
2      int y,  
3      int z)  
4  {  
5      int t1 = x+y;  
6      int t2 = z*48;  
7      int t3 = t1 & 0xFFFF;  
8      int t4 = t2 * t3;  
9      return t4;  
10 }
```

(b) Assembly code

```
      x at %ebp+8, y at %ebp+12, z at %ebp+16  
1      movl    16(%ebp), %eax      z  
2      leal    (%eax,%eax,2), %eax  z*3  
3      sall    $4, %eax           t2 = z*48  
4      movl    12(%ebp), %edx      y  
5      addl    8(%ebp), %edx       t1 = x+y  
6      andl    $65535, %edx        t3 = t1&0xFFFF  
7      imull   %edx, %eax          Return t4 = t2*t3
```

**Figure 3.8** C and assembly code for arithmetic routine body. The stack set-up and completion portions have been omitted.

Does this make sense? Note the usage of `leal` and `sall`, rather than simply using `imull`. Using `imull` isn't wrong - but it's good to be able to see why both approaches work!

# Function Frames

- When a function is called, a section of the stack is set aside for the function.
- Represented by two registers: the base pointer (%ebp) and the stack pointer (%esp).

# **%ebp: base pointer**

- Points to the "beginning" of the function's stack frame.
- Should not change during function execution, unless another function call is made.

# **%ebp: Accessing Arguments**

- Suppose `f()` calls `g(x,y)`. Then, `g` can access its arguments `x,y` via `%ebp`:
  - `x` is at `8(%ebp)`, and `y` is at `12(%ebp)`

# **%ebp**

- What's at 0(%ebp) and 4(%ebp)?
- %ebp points to the saved %ebp, ie the f's base pointer.
- Need to set %ebp to the f's %ebp before returning from g! More on this later.



# **%ebp**

- 4(%ebp) points to the saved return address, i.e. the next instruction to execute after returning from the function.
- The command ret updates the %eip (Instruction Pointer)

# **%esp: Stack Pointer**

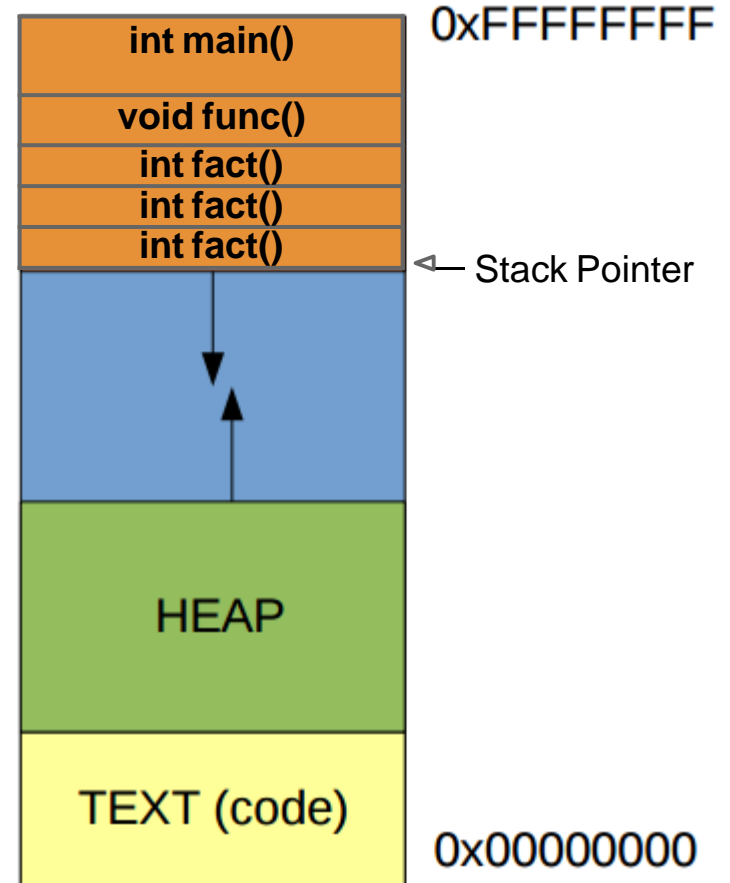
- Points to the "end" of the function frame.
- All of a function's local variables are stored between %ebp and %esp

# **%esp - Static Allocation**

- At the start of a function, we allocate all required storage of local/temp variables by updating %esp

# The Stack

- Contains local variables
- LIFO
- Grows “downward”
- Organized in frames



# The Stack: pushl and popl

- pushl <SRC>
  - subl \$4, %esp
  - movl <SRC> (%esp)
- popl <DST>
  - movl (%esp) <DST>
  - addl \$4 %esp

pushl, popl are convenience commands. Could simply use subl, movl, addl, etc.

# The Stack

Consider the following stack.

What happens when I do:

```
pushl %ebp
```

Addresses  
grow down



0x7fff401c

STACK

**%esp = 0x744401C**

**%ebp = 0x12C**

**%edx = 0x800448B**

# The Stack

Addresses  
grow down



0x7fff401c

0x7fff4018

STACK

0x00000012c

**%esp = 0x7444018**

**%ebp = 0x12C**

**%edx = 0x800448B**

# The Stack

What happens when I  
do:

```
popl %edx
```

Addresses  
grow down



0x7fff401c

0x7fff4018

STACK

0x00000012c

**%esp = 0x7444018**

**%ebp = 0x12C**

**%edx = 0x800448B**



# The Stack

Addresses  
grow down



STACK

0x7fff401c

**%esp = 0x744401C**

**%ebp = 0x12C**

**%edx = 0x12C**

# The Stack

- How do we create frame abstractions for procedures?

Addresses  
grow down



0x7fff401c

STACK

**%esp = 0x744401C**  
**%ebp = 0x12C**  
**%edx = 0x12C**

# Stack Frames

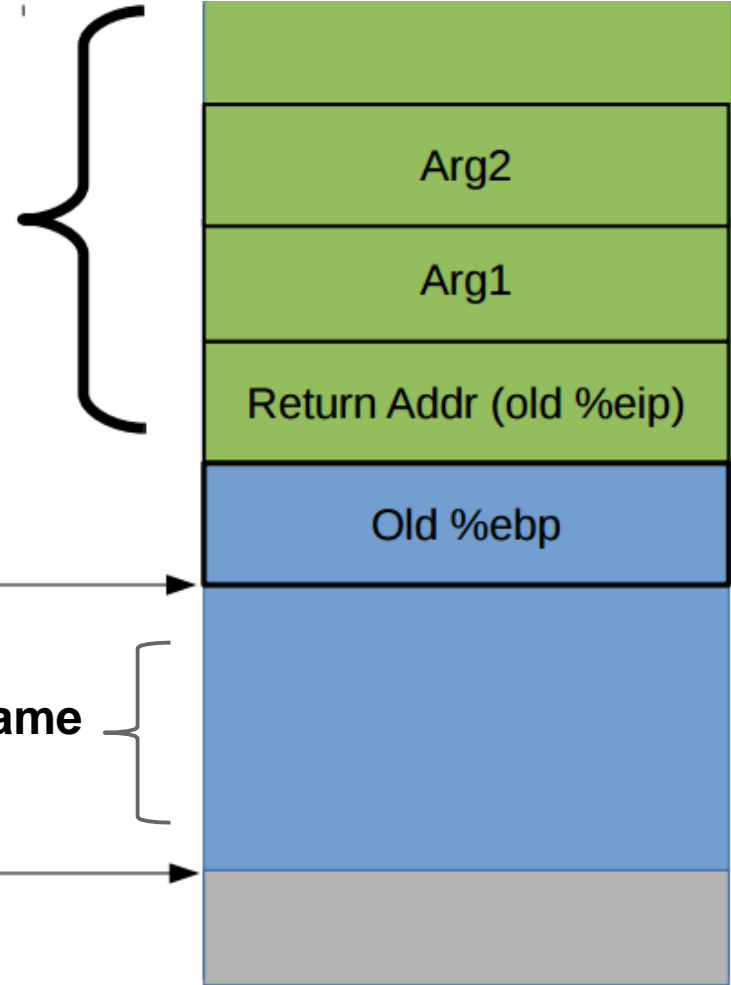
- `%ebp`
  - base pointer
  - bottom of frame
- `%eip`
  - instruction pointer

Caller Frame

Current Frame

`%esp`

`%ebp`



# x86 Calling Conventions

- **Caller** saved registers
  - pushed to the stack before function call is made
  - restored after the callee exits
  - %eax, %ecx, %edx
- **Callee** saved registers
  - pushed to the stack at the start of the function
  - restored before the callee exits
  - %ebx, %edi, %esi
  - return value in %eax

# Even More Assembly

- `call <label>`
  - `pushl %eip`
  - `movl <label (function address)> %eip`
- `ret`
  - `popl %eip`

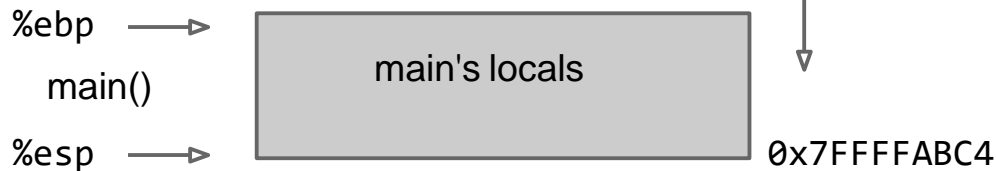
# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x);  
    return foo + 5;  
}  
  
int g(int num) {  
    return num+10;  
}
```

```
In main():  
    f();
```

# Example: Function Call

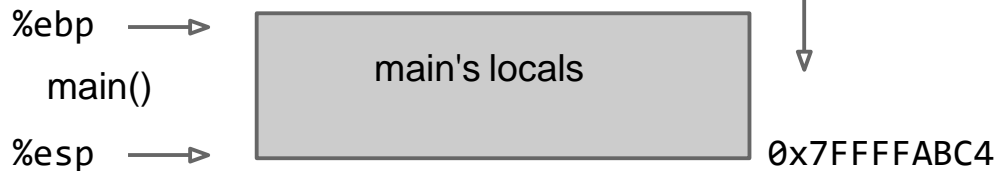
```
int f() {  
    int x = 42;  
    int foo = g(x);  
    return foo + 5;  
}  
  
int g(int num) {  
    return num+10;  
}
```



In `main()`, about to call `f()`.

# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x);  
    return foo + 5;  
}  
  
int g(int num) {  
    return num+10;  
}
```

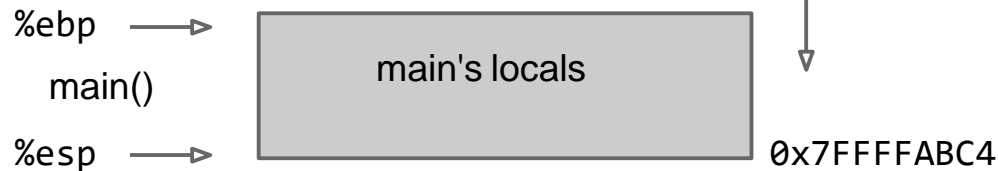


In `main()`, about to call `f()`.  
1. Prepare arguments to `f()`.



# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x);  
    return foo + 5;  
}  
  
int g(int num) {  
    return num+10;  
}
```

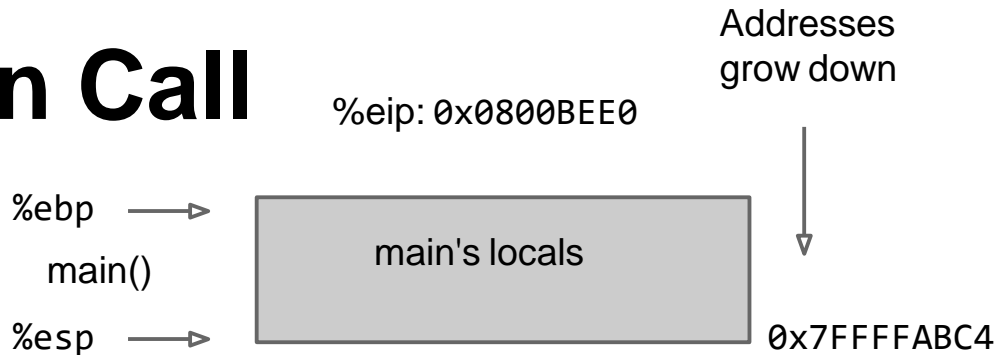


In `main()`, about to call `f()`.  
1. Prepare arguments to `f()`.

**No arguments!**

# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x);  
    return foo + 5;  
}  
  
int g(int num) {  
    return num+10;  
}
```



In `main()`, about to call `f()`.  
1. Prepare arguments to `f()`.  
2. Call `f()`

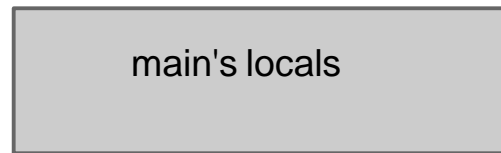
# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x);  
    return foo + 5;  
}  
  
int g(int num) {  
    return num+10;  
}
```

%ebp →

main()

%esp →



%eip: 0x0800BEE0

Addresses  
grow down



0x7FFFFABC4

Assume instruction after call to g(x) is at  
location: 0x0800BEE4

Assume first instruction of g is at location:  
0x0800F00D

In main(), about to call f().

1. Prepare arguments to f().
2. Call f()

call f

# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x);  
    return foo + 5;  
}  
  
int g(int num) {  
    return num+10;  
}
```



In main(), about to call f().

1. Prepare arguments to f().
2. Call f()

call f

# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x);  
    return foo + 5;  
}  
  
int g(int num) {  
    return num+10;  
}
```



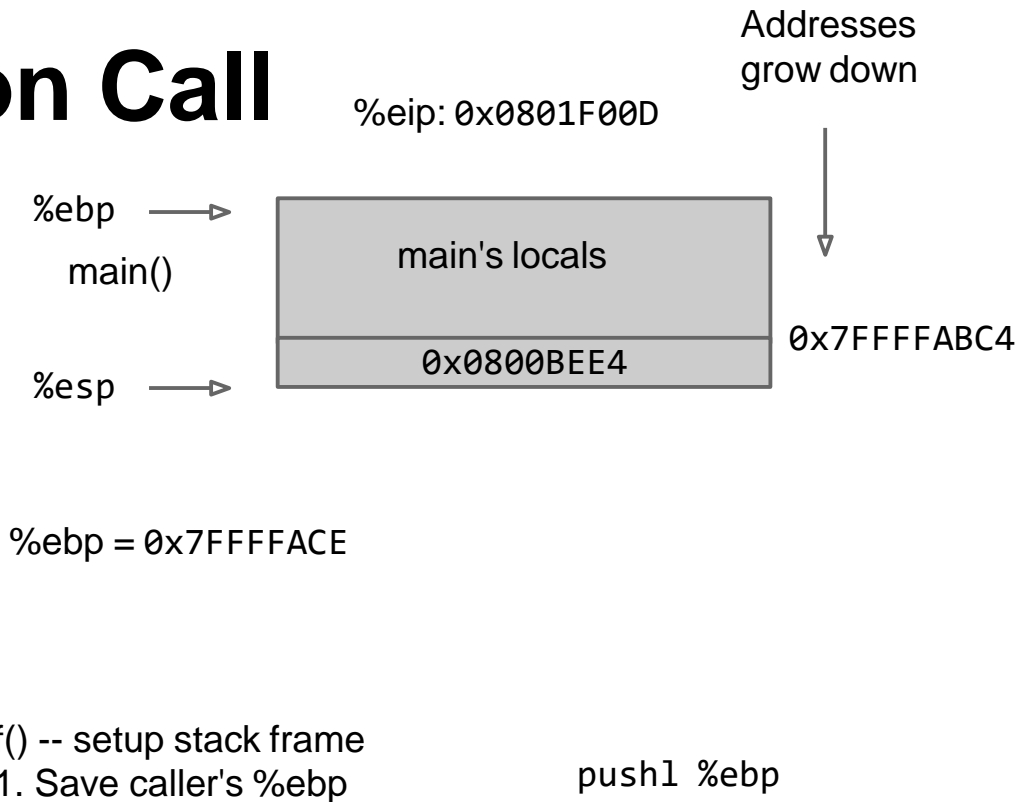
In main(), about to call f().

1. Prepare arguments to f().
2. Call f()
3. f() is now "in control"

call f

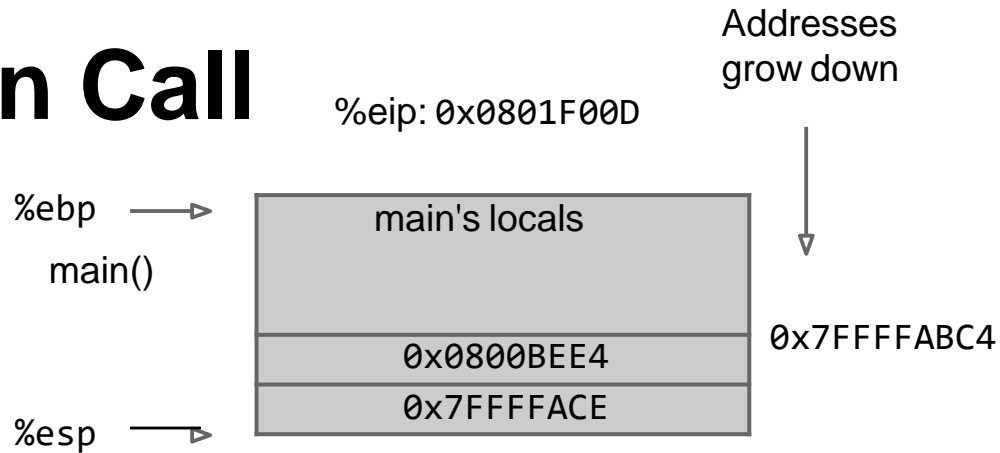
# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x);  
    return foo + 5;  
}  
  
int g(int num) {  
    return num+10;  
}
```



# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x);  
    return foo + 5;  
}  
  
int g(int num) {  
    return num+10;  
}
```

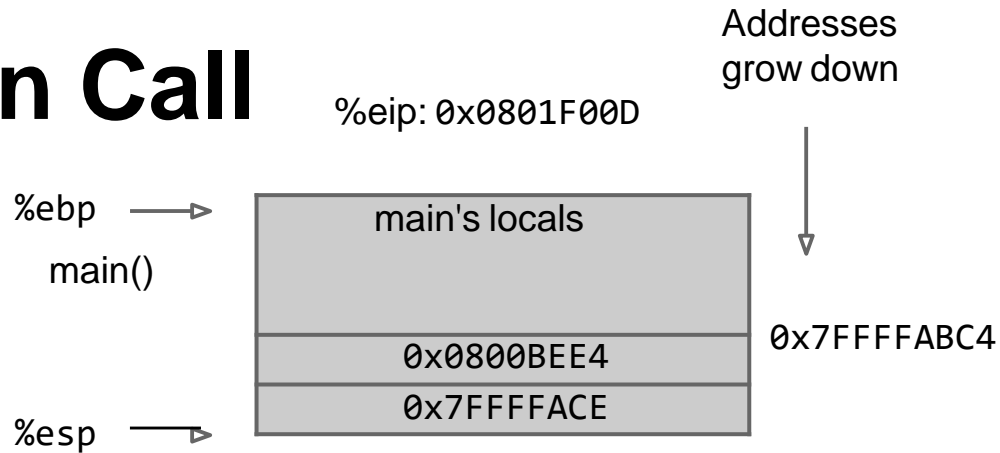


f() -- setup stack frame  
1. Save caller's `%ebp`

`pushl %ebp`

# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x);  
    return foo + 5;  
}  
  
int g(int num) {  
    return num+10;  
}
```



f() -- setup stack frame

1. Save caller's %ebp
2. Update %ebp to point to \*my\* frame base.

movl %esp %ebp



# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x);  
    return foo + 5;  
}  
  
int g(int num) {  
    return num+10;  
}
```



f() -- setup stack frame

1. Save caller's %ebp
2. Update %ebp to point to \*my\* frame base.

movl %esp %ebp

# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x);  
    return foo + 5;  
}  
  
int g(int num) {  
    return num+10;  
}
```



f() -- setup stack frame

1. Save caller's %ebp

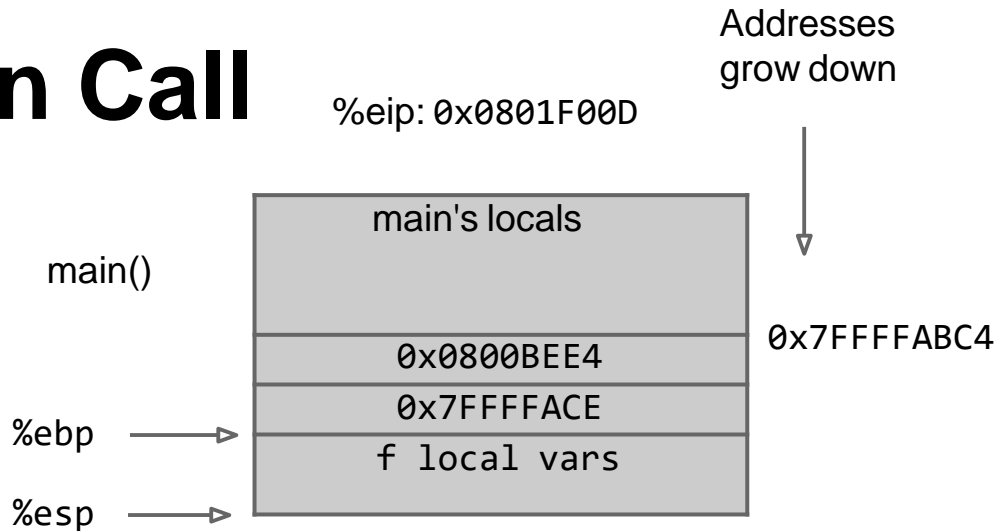
2. Update %ebp to point to \*my\*  
frame base.

3. Allocate space for local  
variables

```
subl $8 %esp
```

# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x);  
    return foo + 5;  
}  
  
int g(int num) {  
    return num+10;  
}
```



*Note: Depending on how clever the compiler is, it may not allocate stack space for \*all\* local vars if it can do it with registers.*

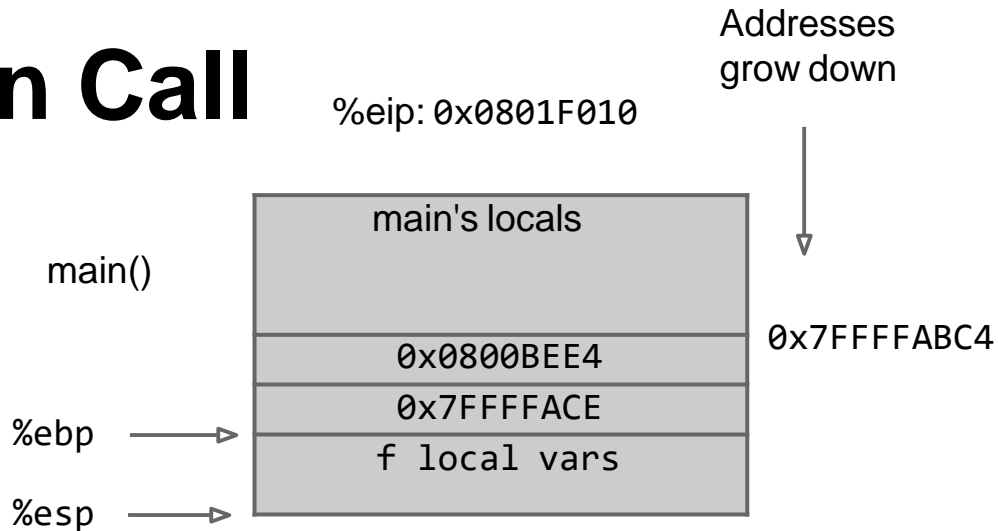
f() -- setup stack frame

1. Save caller's %ebp
2. Update %ebp to point to \*my\* frame base.
3. Allocate space for local variables

```
subl $8 %esp
```

# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x);  
    return foo + 5;  
}  
  
int g(int num) {  
    return num+10;  
}
```

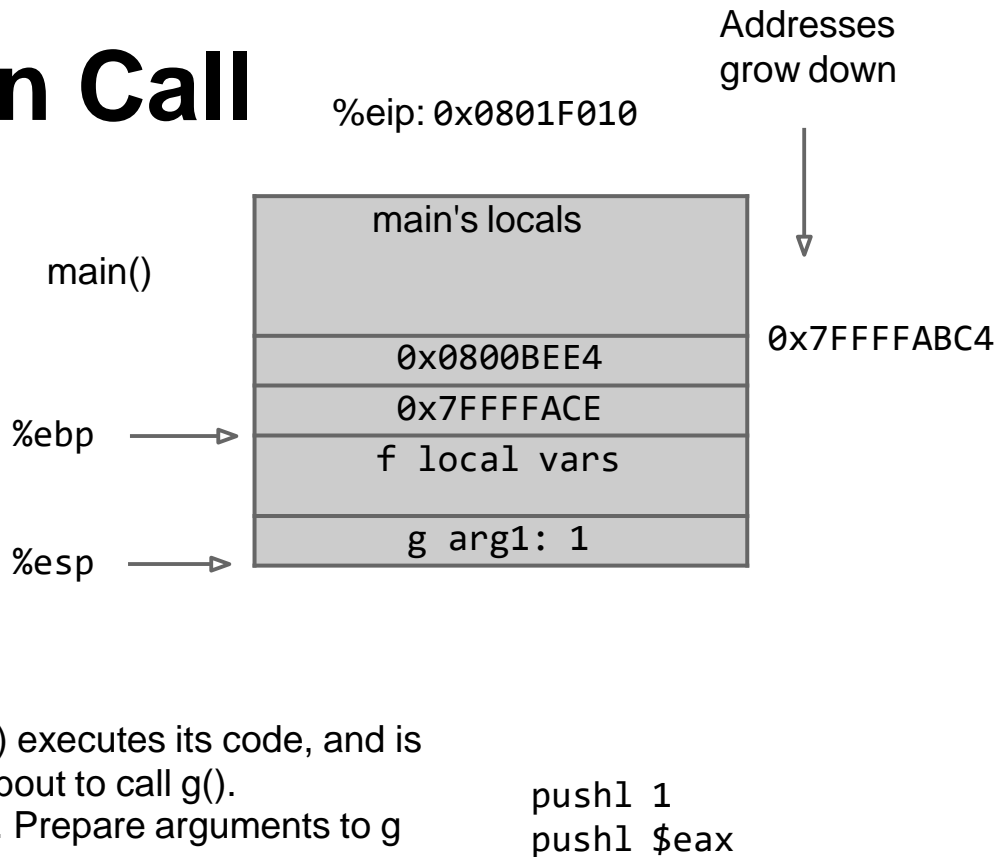


f() executes its code, and is about to call g().  
1. Prepare arguments to g

pushl \$eax

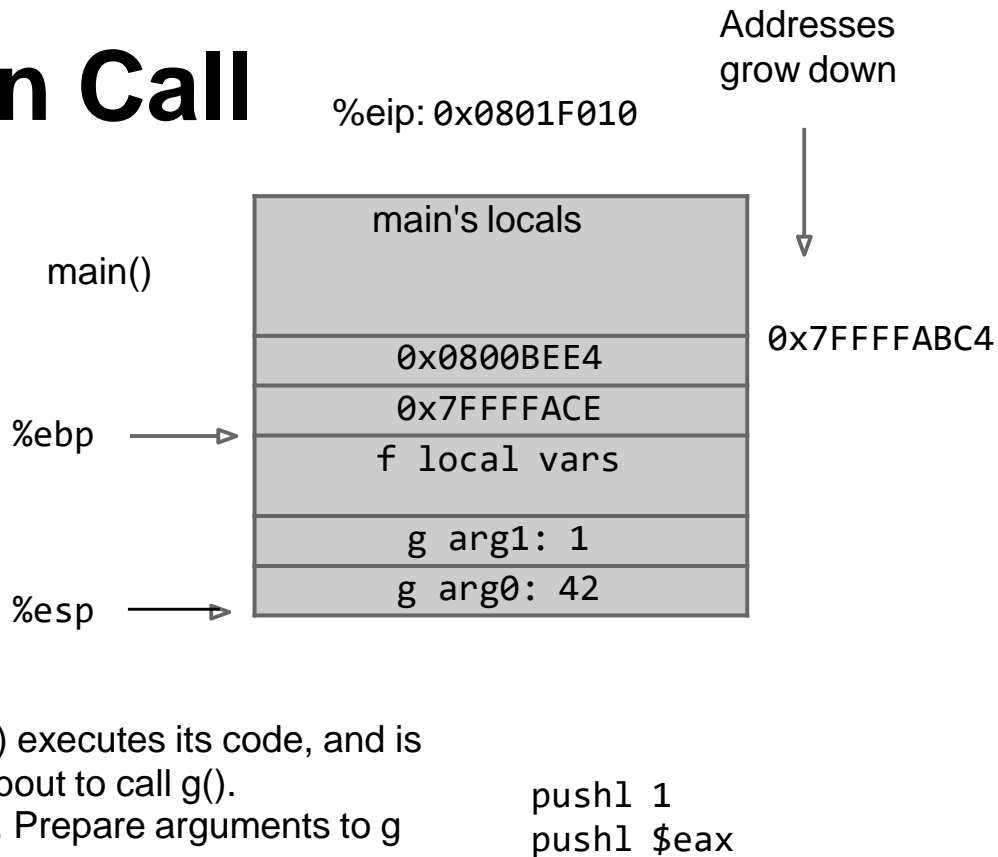
# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x,1);  
    return foo + 5;  
}  
  
int g(int n, int a) {  
    return num+10;  
}
```



# Example: Function Call

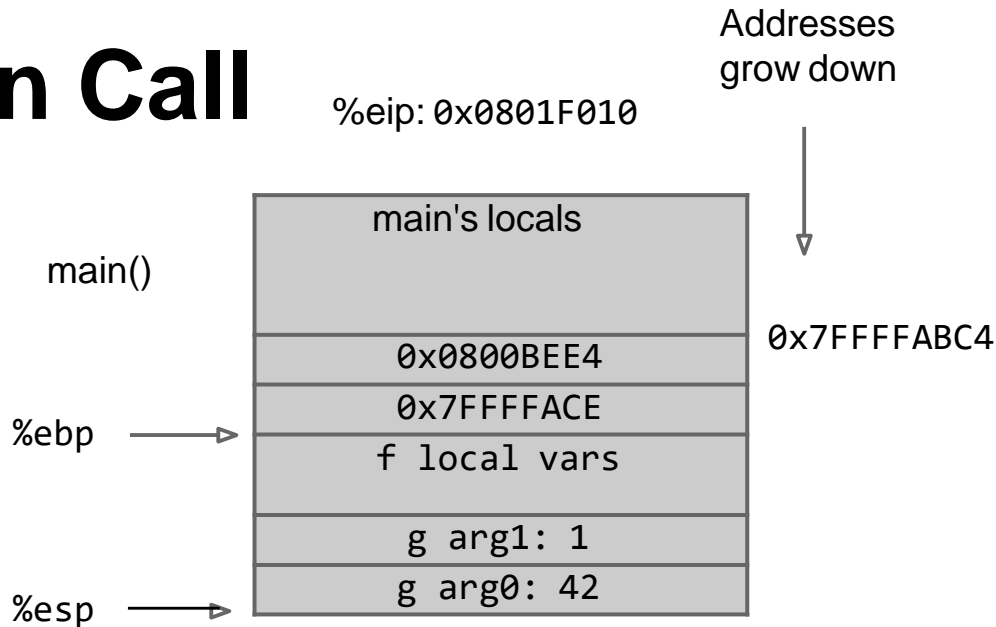
```
int f() {  
    int x = 42;  
    int foo = g(x,1);  
    return foo + 5;  
}  
  
int g(int n, int a) {  
    return num+10;  
}
```



# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x,1);  
    return foo + 5;  
}  
int g(int n, int a) {  
    return num+10;  
}
```

Assume f's next instruction  
is at 0x0801F014



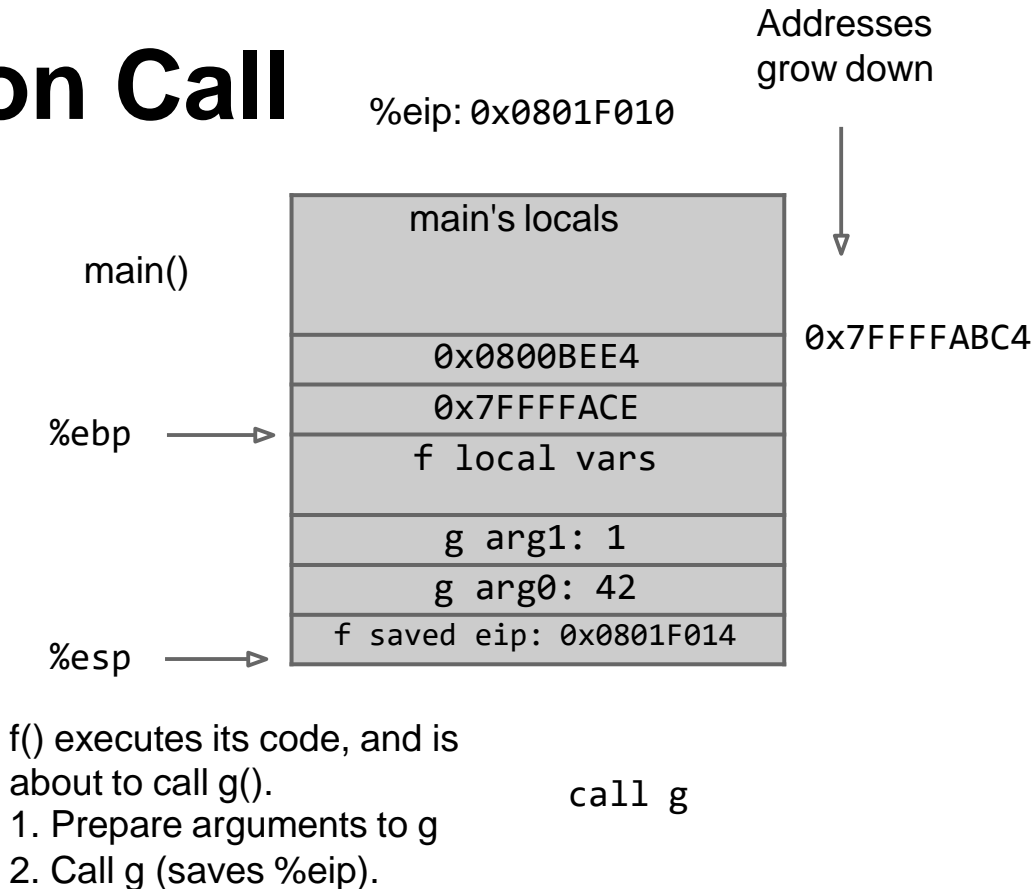
f() executes its code, and is  
about to call g().  
1. Prepare arguments to g  
2. Call g (saves %eip).

call g

# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x,1);  
    return foo + 5;  
}  
int g(int n, int a) {  
    return num+10;  
}
```

Assume f's next instruction  
is at 0x0801F014

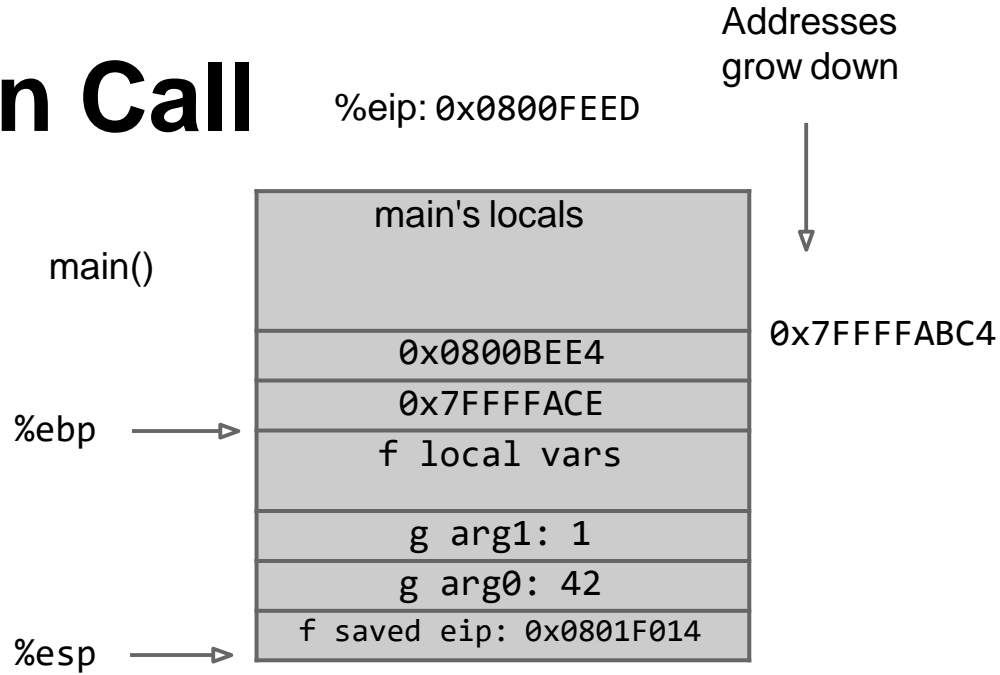




# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x,1);  
    return foo + 5;  
}  
int g(int n, int a) {  
    return num+10;  
}
```

Assume f's next instruction  
is at 0x0801F014



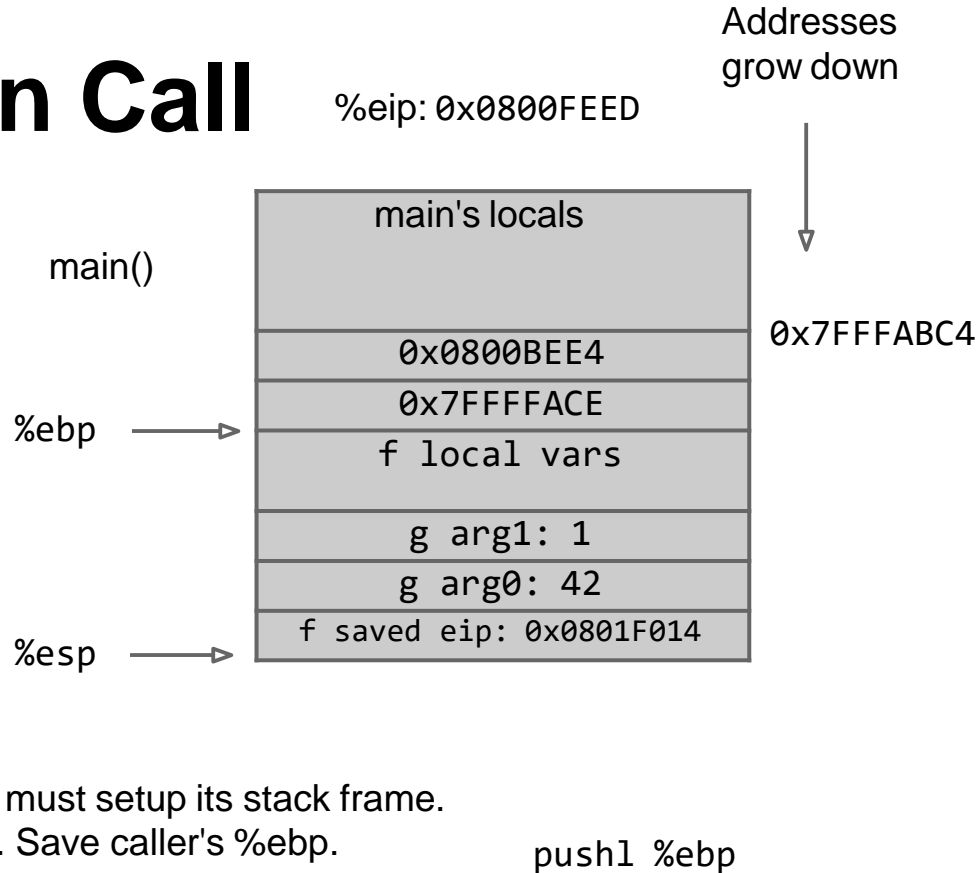
f() executes its code, and is  
about to call g().

1. Prepare arguments to g
2. Call g (saves %eip).
3. g is in control now!

call g

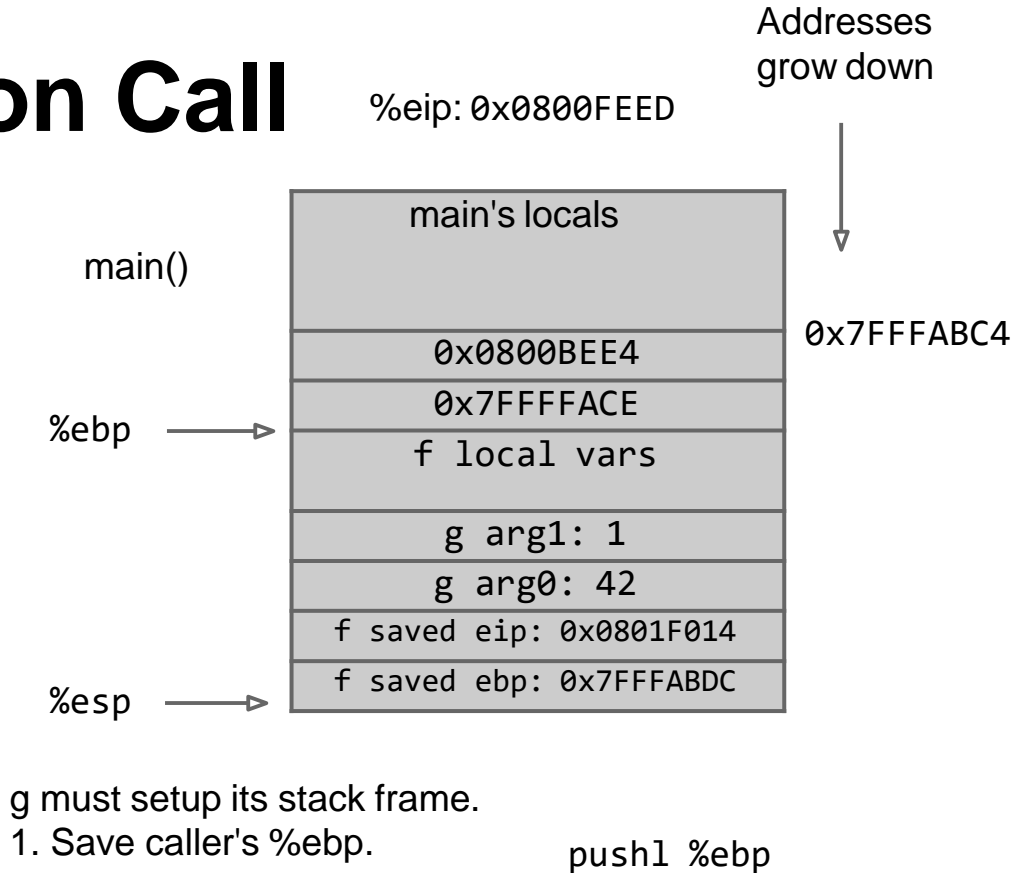
# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x,1);  
    return foo + 5;  
}  
int g(int n, int a) {  
    return num+10;  
}
```



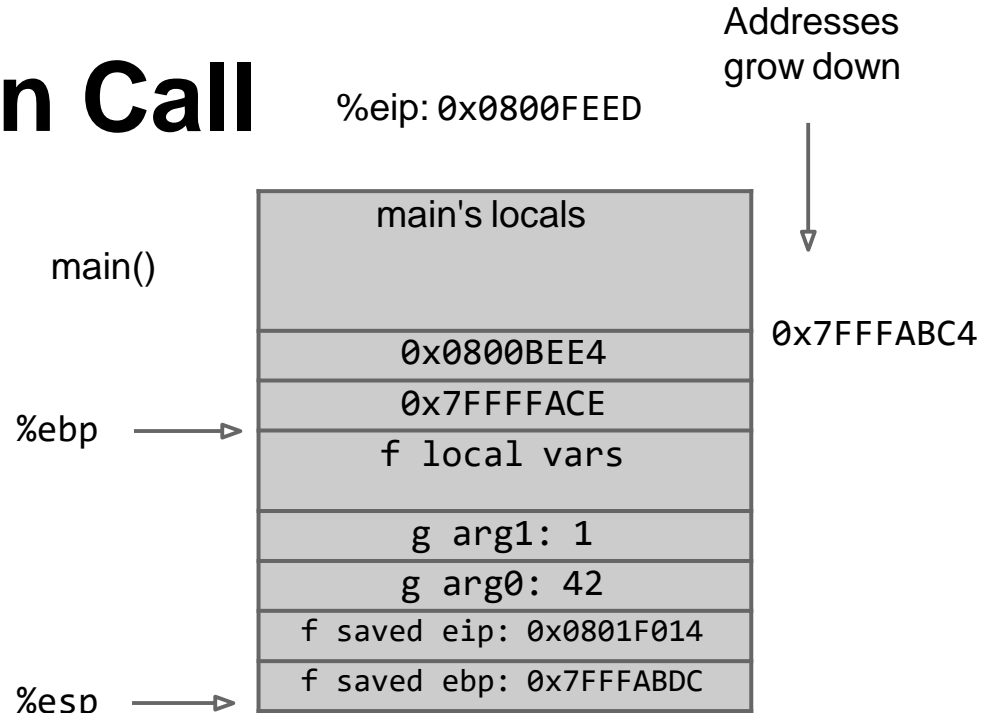
# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x,1);  
    return foo + 5;  
}  
int g(int n, int a) {  
    return num+10;  
}
```



# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x,1);  
    return foo + 5;  
}  
int g(int n, int a) {  
    return num+10;  
}
```



g must setup its stack frame.

1. Save caller's %ebp.
2. Update %ebp to point to \*my\* frame base.

movl %esp %ebp

# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x,1);  
    return foo + 5;  
}  
int g(int n, int a) {  
    return num+10;  
}
```

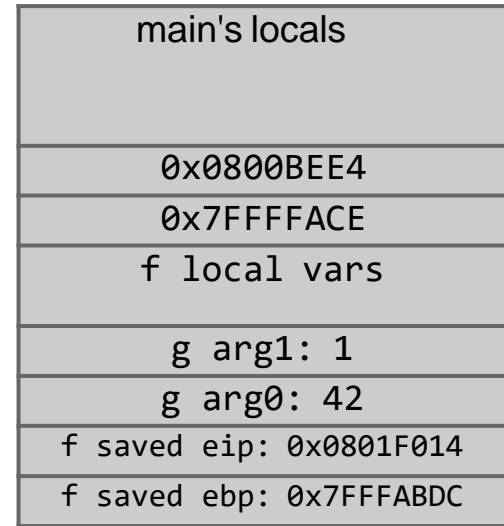
main()

%eip: 0x0800FEED

Addresses  
grow down



0x7FFFABC4



%ebp

%esp



g must setup its stack frame.

1. Save caller's %ebp.
2. Update %ebp to point to \*my\* frame base.

movl %esp %ebp

# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x,1);  
    return foo + 5;  
}  
int g(int n, int a) {  
    return num+10;  
}
```

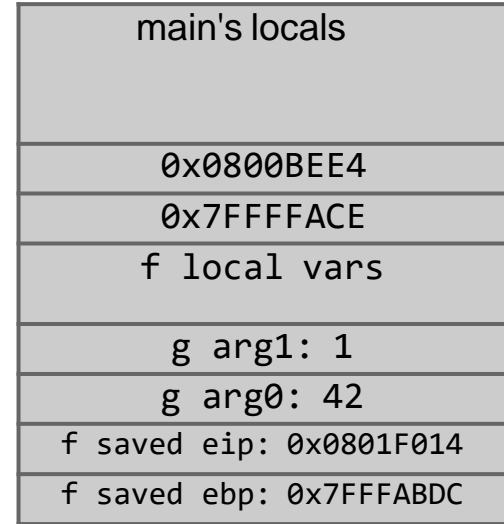
main()

%eip: 0x0800FEED

Addresses  
grow down



0x7FFFABC4



%ebp

%esp



g must setup its stack frame.

1. Save caller's %ebp.
2. Update %ebp to point to \*my\* frame base.
3. Allocate space for local vars.

# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x,1);  
    return foo + 5;  
}  
  
int g(int n, int a) {  
    return n+10;  
}
```

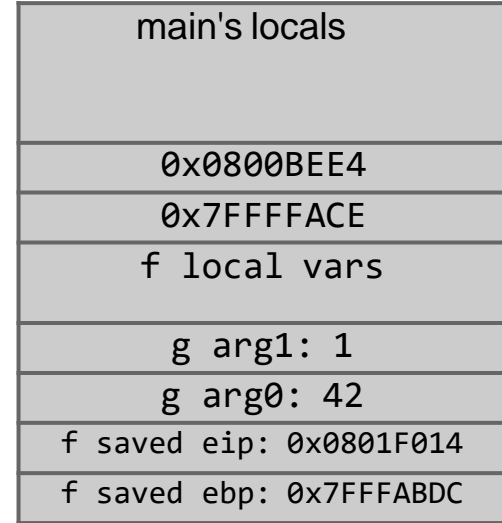
main()

%eip: 0x0800FEED

Addresses  
grow down



0x7FFFABC4



%ebp

%esp



g must setup its stack frame.

1. Save caller's %ebp.
2. Update %ebp to point to \*my\* frame base.
3. Allocate space for local vars.

**No local vars!**

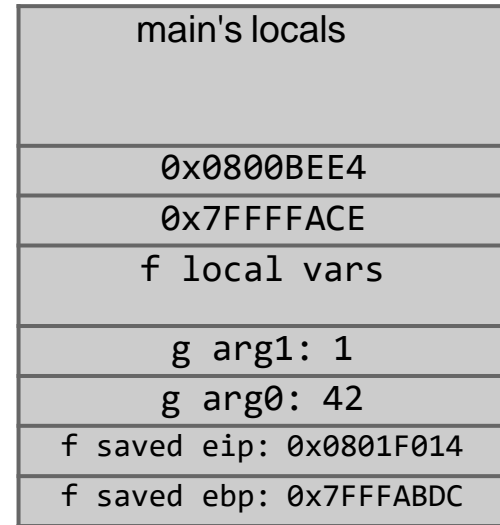
# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x,1);  
    return foo + 5;  
}  
int g(int n, int a) {  
    return n+10;  
}
```

main()

%eip: 0x0800FEED

Addresses  
grow down



%ebp

%esp



g() finishes, now must return.

1. Set %esp to caller's original %esp.

2. Set %ebp to caller's original %ebp.

**This is %ebp!**

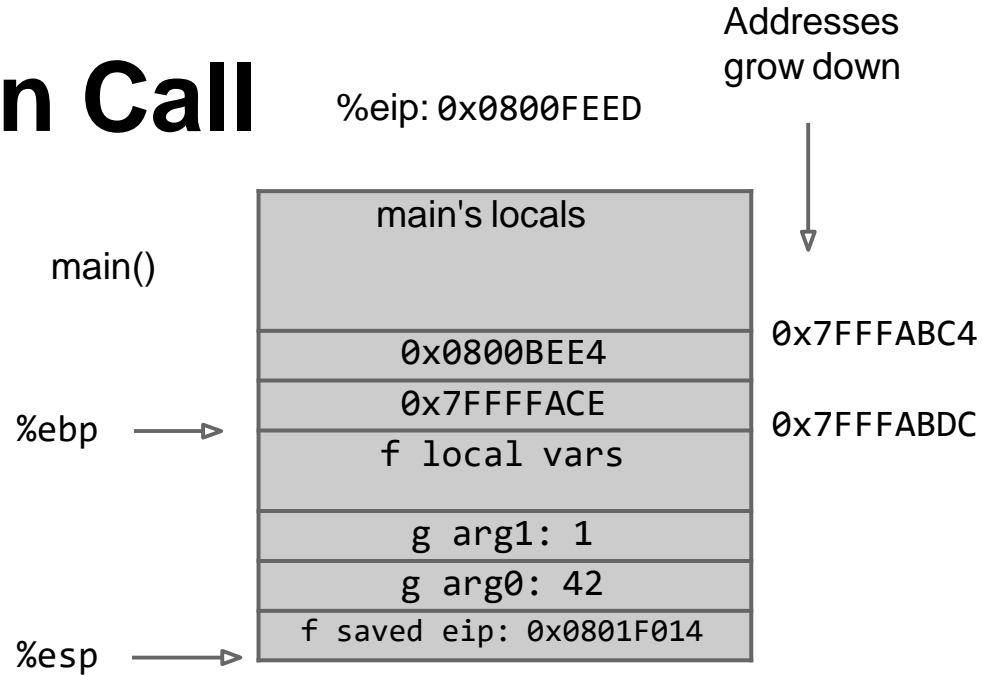
movl %ebp %esp

popl %ebp



# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x,1);  
    return foo + 5;  
}  
int g(int n, int a) {  
    return n+10;  
}
```



g() finishes, now must return.

1. Set %esp to caller's original %esp.

2. Set %ebp to caller's original %ebp.

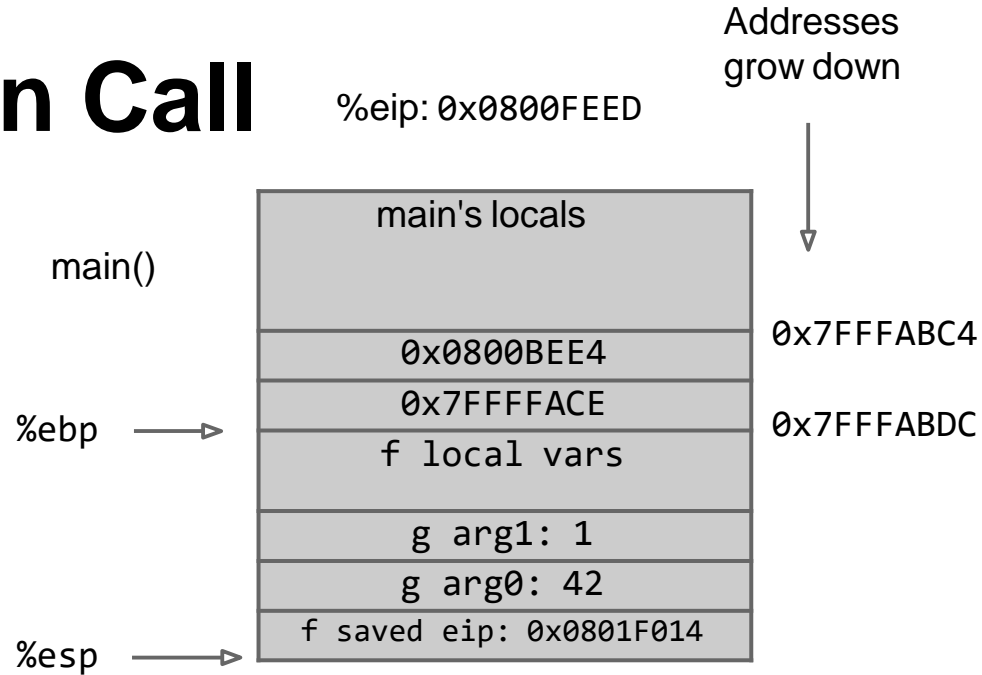
**This is %ebp!**

movl %ebp %esp

popl %ebp

# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x,1);  
    return foo + 5;  
}  
  
int g(int n, int a) {  
    return n+10;  
}
```



g() finishes, now must return.

1. Set %esp to caller's original %esp.
2. Set %ebp to caller's original %ebp.
3. Return to caller.

***Pops top of stack,  
places value into  
%eip.***

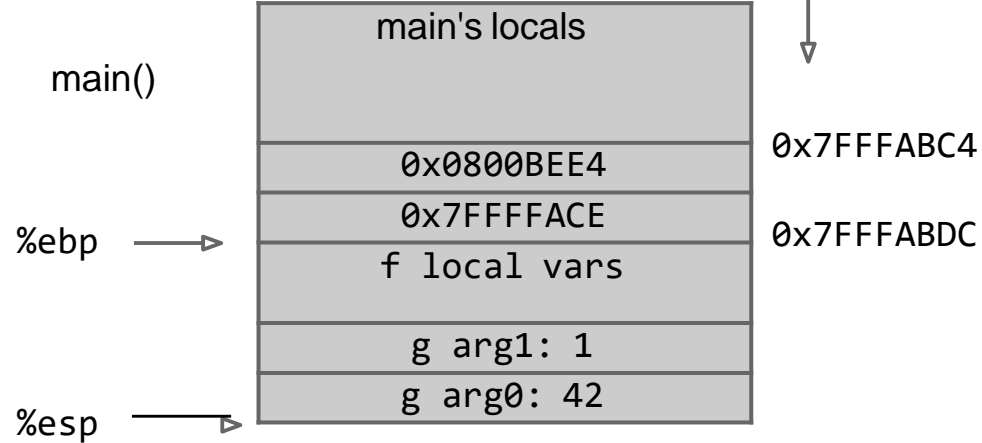
ret

# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x,1);  
    return foo + 5;  
}  
  
int g(int n, int a) {  
    return n+10;  
}
```

f in control now!  
%eip: 0x0801F014

Addresses  
grow down



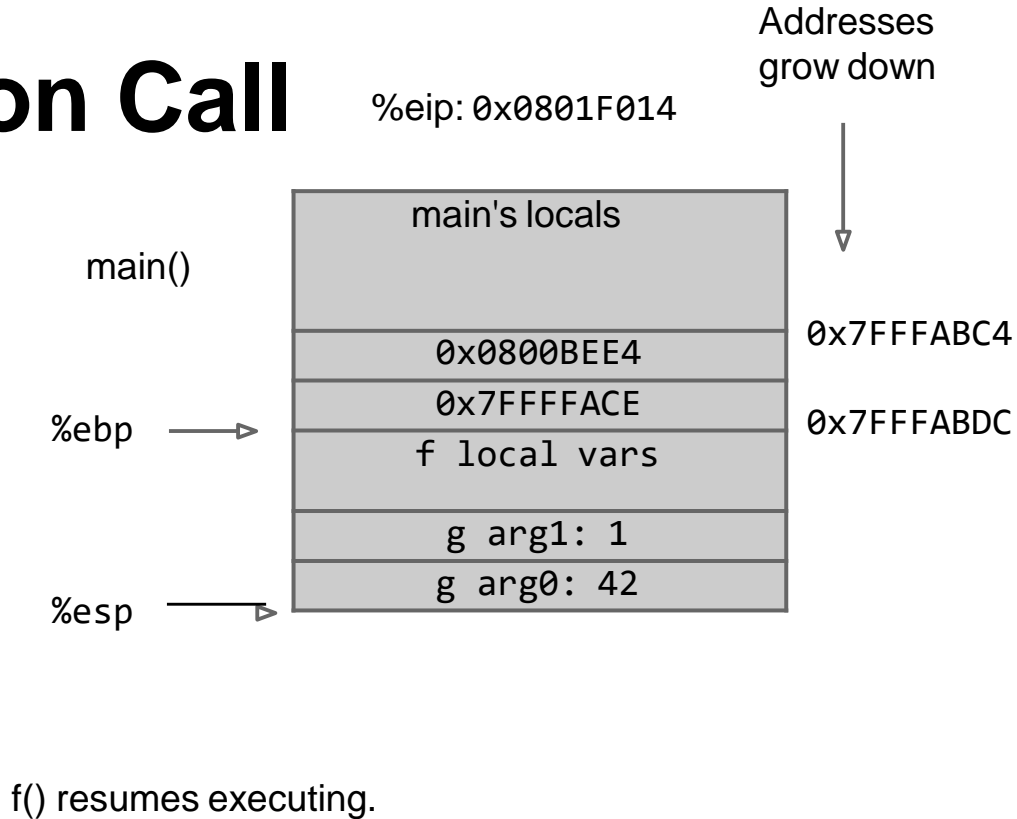
g() finishes, now must return.  
1. Set %esp to caller's original % esp.  
2. Set %ebp to caller's original % ebp.  
3. Return to caller.

***Pops top of stack,  
places value into  
%eip.***

ret

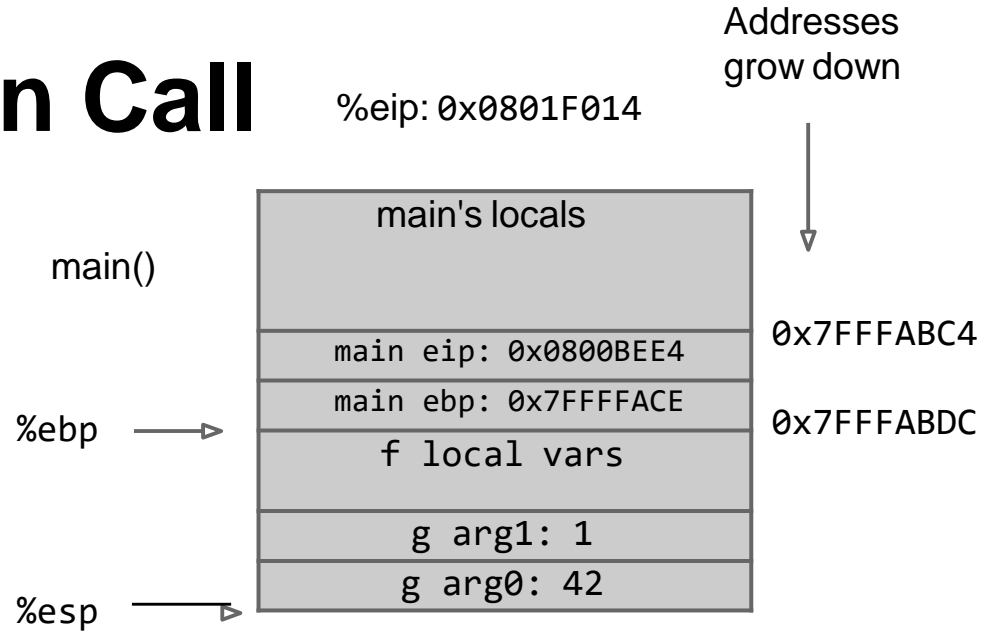
# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x,1);  
    return foo + 5;  
}  
  
int g(int n, int a) {  
    return n+10;  
}
```



# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x,1);  
    return foo + 5;  
}  
  
int g(int n, int a) {  
    return n+10;  
}
```

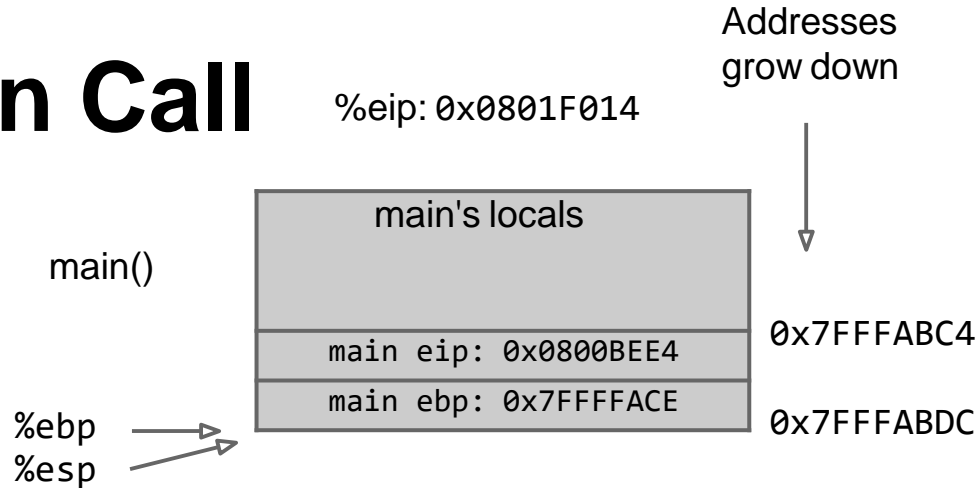


f() is ready to return.

1. Update %esp to caller's %esp.      `movl %ebp %esp`

# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x,1);  
    return foo + 5;  
}  
int g(int n, int a) {  
    return n+10;  
}
```



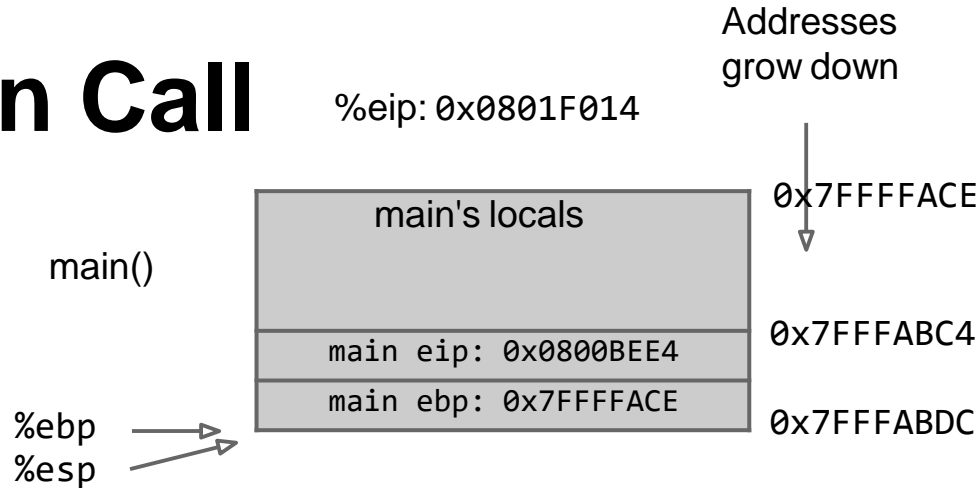
`f()` is ready to return.

1. Update `%esp` to caller's `%esp`.

`movl %ebp, %esp`

# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x,1);  
    return foo + 5;  
}  
int g(int n, int a) {  
    return n+10;  
}
```



f() is ready to return.

1. Update %esp to caller's %esp.
2. Update %ebp to caller's %ebp.

popl %ebp

# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x,1);  
    return foo + 5;  
}  
int g(int n, int a) {  
    return n+10;  
}
```



f() is ready to return.

1. Update %esp to caller's %esp.
2. Update %ebp to caller's %ebp.

popl %ebp



# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x,1);  
    return foo + 5;  
}  
int g(int n, int a) {  
    return n+10;  
}
```



`f()` is ready to return.

1. Update `%esp` to caller's `%esp`.
2. Update `%ebp` to caller's `%ebp`.
3. Return control to caller.

`ret`

# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x,1);  
    return foo + 5;  
}  
int g(int n, int a) {  
    return n+10;  
}
```



main() in control!  
`%eip: 0x0800BEE4`

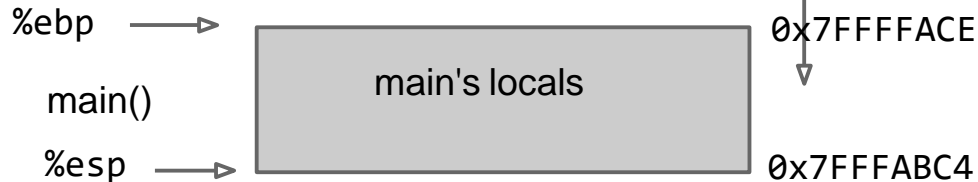
`f()` is ready to return.

1. Update `%esp` to caller's `%esp`.
2. Update `%ebp` to caller's `%ebp`.
3. Return control to caller.

`ret`

# Example: Function Call

```
int f() {  
    int x = 42;  
    int foo = g(x,1);  
    return foo + 5;  
}  
  
int g(int n, int a) {  
    return n+10;  
}
```



main() in control!  
`%eip: 0x0800BEE4`

Addresses  
grow down

main() resumes executing where  
it left off, and finishes its  
awesome computation.

.main:

```
...  
call f  
# here now  
...
```

# Conditional Jumps

`je, jz` -- jump if equal/zero

`jne, jnz` -- jump if not-equal/not-zero

`j1, j1e` -- jump if less than/less-than-or-equal

`jg, jge` -- jump if greater than/greater-than-or-equal

Several more jump types (ie overflow, sign, parity, etc.).

# cmp

Use cmpl, testl to use the conditional jump!

```
cmpl %eax %edx
```

```
jge .L2
```

# testl

```
testl %eax, %eax
```

```
jz zeroLabel; jump if %eax is zero
```

```
js negLabel ; jump if EAX is negative
```

```
jns posLabel ; jump if EAX is positive
```

Quick way to check if a register is 0, negative, or positive.

# x86-64 Stack Convention

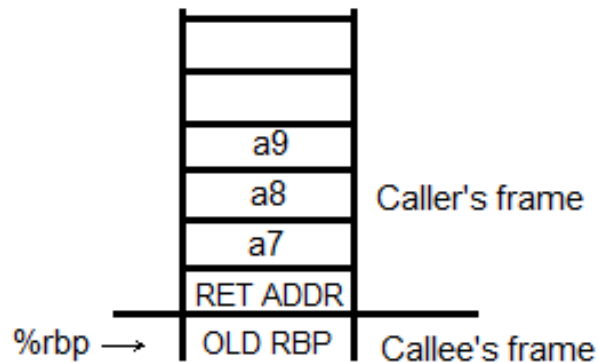
- What would the stack look like if we passed
- more than 6 arguments
- into the function?

Ex:

```
int func(long a1, long a2,  
long a3, long a4, long a5,  
long a6, long a7, long a8,  
long a9)
```

# x86-64 Stack Convention

- `int func (long a1, long a2, long a3, long a4, long a5, long a6, long a7, long a8, long a9)`
- The first six arguments will be stored in registers.
- The following arguments will be placed on the stack by the caller.
- The earlier the argument in the list, the closer to `%rbp`.
- In this example, `a7` is accessible to the callee via `0x10(%rbp)`, `a8` is accessible via `0x18(%rbp)`, etc.



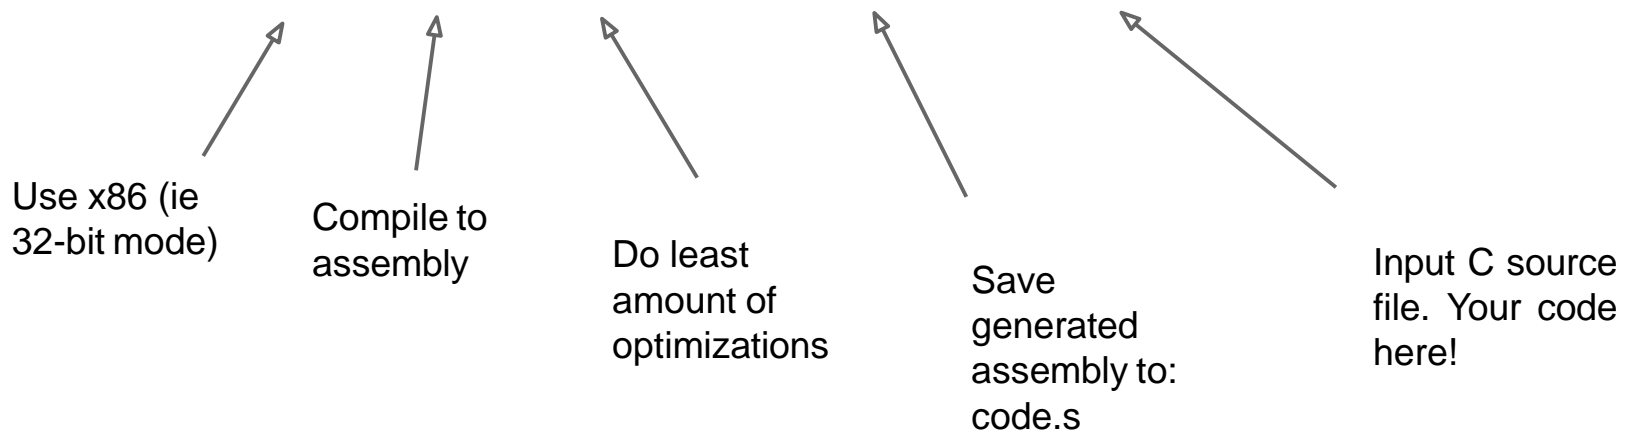


# Compiling at Home

Try creating assembly output yourself!

```
$ gcc -m32 -S -O0 -o code.s code.c
```

Use x86 (ie  
32-bit mode)



The diagram consists of five arrows pointing upwards from explanatory text blocks to specific flags in the command line. The first arrow points from 'Use x86 (ie 32-bit mode)' to '-m32'. The second arrow points from 'Compile to assembly' to '-S'. The third arrow points from 'Do least amount of optimizations' to '-O0'. The fourth arrow points from 'Save generated assembly to: code.s' to '-o code.s'. The fifth arrow points from 'Input C source file. Your code here!' to 'code.c'.

Compile to  
assembly

Do least  
amount of  
optimizations

Save  
generated  
assembly to:  
code.s

Input C source  
file. Your code  
here!

# Compiling at Home

Add this flag to disable weird lines with `.cfi_` junk:

```
$ gcc -m32 -S -O0 -fno-asynchronous-unwind-tables -o code.s code.c
```

# Compiling at Home

Full pipeline to compile .c code -> .s -> executable.

# Compile: generates assembly from c code

```
gcc -S -m32 -O0 -fno-asynchronous-unwind-tables -o printint.s  
printint.c
```

# Assemble: generates object file from

```
assembly gcc -c -m32 -o printint.o printint.s
```

# Linker: generates executable from object

```
file gcc -m32 -o printint printint.o
```

Use last two commands to create executables of your own x86 code!

# Structs

```
struct s {  
    char c1;  
    int i;  
    char c2;  
    int j;  
};
```

- What's the problem with this struct?

# Structs

```
struct s {  
    char c1;  
    int i;  
    char c2;  
    int j;  
};
```

- Say an instance of the struct begins at 0x10. Then c1 is at address 0x10. However, 'i' cannot be at address 0x11 (it needs to be 4-aligned). As a result, we need 3 bytes of padding.

# Structs

- This is a waste of space! There will be 3 bytes of padding after c1 and 3 bytes of padding after c2, meaning that this struct will take up 16 bytes when really it only needs 10.

# Structs

- Two common struct ordering guidelines (which could be at odds):
  1. Place the most commonly used data type first.
  2. Place the elements in descending order of size (ie largest first)
- Why?

# Structs

- 1.
- Memory references are expensive (ex. (%eax))... but memory references with an offset are more expensive (ex. 8(%eax))
- Chances are, you'll be referring to the struct by a pointer to the beginning of the struct, which means that dereferencing the pointer without an offset will point to the first element.



# Structs

- 2.
- If the elements with larger sizes are first, that means there will be less of a need for padding.
- For example, consider struct s, except with the first two elements swapped:

```
struct s {  
    int i;  
    char c1;  
    char c2;  
    int j;  
};
```

# Structs

- 2.
- ```
struct s {  
    int i;  
    char c1;  
    char c2;  
    int j;  
};
```
- Now, we need 2 bytes of padding between c2 and j for a total of 12 bytes.

# Structs

- Because each internal element must follow their own alignment rules, the alignment of the struct must be equal to the strictest of the elements within a struct.
- But wait...

# Structs

- Consider:

```
struct s {  
    char c;  
    int i;  
};
```

- Because int i is aligned by 4, instances of struct s must be aligned by 4.
- There must also be 3 bytes of padding between c and i, meaning a total size of 8.

# Structs

- Thus, a possible placement of (struct s s1) where s1.c = 0xFF and s1.i = 0x33221100 is the following:

|          |      |      |      |      |      |      |      |      |
|----------|------|------|------|------|------|------|------|------|
| Address: | 0x10 | 0x11 | 0x12 | 0x13 | 0x14 | 0x15 | 0x16 | 0x17 |
| Value:   | 0xFF | 0xFF | 0xFF | 0xFF | 0x00 | 0x11 | 0x22 | 0x33 |

- Where s begins at 0x10.
- This is how we meet the alignment requirements of each individual item

# Structs

- But wait... what if you had the following code:

```
struct T
{
    struct s foo;
    char c;
} t;
```

- If t began at 0x10:
  - t.foo.i : 0x10 → 0x13
  - t.foo.c : 0x14
- But wait, t.c doesn't have to start at 0x18. It can be at 0x15 and all of the rules will be followed, right?

# Structs

- But does it? If you try this:

```
int main()
{
    struct T test;
    printf("%p\n", &test.foo.i);
    printf("%p\n", &test.foo.c);
    printf("%p\n", &test.c);
}
```

- Output:
  - 0x7ffdca140b40
  - 0x7ffdca140b44
  - 0x7ffdca140b48
- Nope, looks like sizeof(s) is really 8 bytes.

# Unions

- Like structs except all of the values begin at the same address.
- union s {
  - short s;
  - char c;
  - };
- This means that in a union that contains several values, only one of them is likely to be meaningful and assigning one term a value will trample other terms.



# Unions

- union s {
- short s;
- char c;
- };
- union s foo;
- Say foo begins at 0x10.
- foo.s will be located in addresses 0x10 and 0x11
- foo.c will be located in address 0x10.

# Unions

- union s {
- short s;
- char c;
- };
- union s foo;
- foo.s = 0xFFFF;
- foo.c = 0;
- printf(“%hx\n”, foo.s) => FF00

# Midterm

- On Feb 7
- Covers the lecture material and relevant topics from book
- Covers the assignments and labs