

# Some Review

# Test Logistics

- Date: Monday, March 20th
- Time: 3-6pm (but likely many will finish sooner)
- Place: Kinsey Science Teaching Pavilion: 1220B
  - (Ignore the other room)

# Review

## ■ Data Representations:

- Integers, Bit Manipulation, Floats, Casting & Data Types

## ■ Machine Level Programming:

- Instruction Types, Control Flow (conditionals, procedure calls), Data Types (struct, union), Buffer Overflow

## ■ Memory Hierarchy:

- Principle of caching, locality (temporal, spatial), Memory Hierarchy(disks, memory, cache, register), Virtual memory (TLBs)

## ■ Low-level Program Optimizations:

- Analysis: Calculating latency, CPE, cache miss rates
- Optimizations for cache (blocking, loop ordering), function inlining, unrolling, strength reduction

## ■ Parallelism:

- Threading (pthreads), race conditions, semaphores, deadlocks.

## ■ System Guts:

- Linking, System Calls, Traps/Exceptions, Forking/Loading

# What's the test going to look like?

## ■ Many multiple choice / short answer questions

- Cover miscellaneous topics – there is no material that is out of bounds. (though I will try to not be mean : )
- Naming which part of the system is responsible for what

## ■ Straightforward Questions:

- Given a piece of C code, identify where those variables are stored, and other properties (shared, private)
- Analyze a given piece of code for performance, explain which optimizations could be applied
- Virtual Memory + Caching

# Where's that Memory?

*main.c*

```
int sum(int *a, int n);

int array[2] = {1, 2};

int main()
{
    int val1 = sum(array, 2);
    int * dyn_arr = (int *) malloc(4*2);
    int val2 = sum(array, 2);
    return val1 + val2;
}
```

.data

.heap

Stored in %rsi,  
%rdi

Stored in

.text

*sum.c*

```
int sum(int *a, int n)
{
    int i, s = 0;
    struct_t n;

    for (i = 0; i < n; i++) {
        s += a[i];
    }
    return s;
}
```

Stored on  
stack

:)

Probably a register

%rax

# Code Optimizations

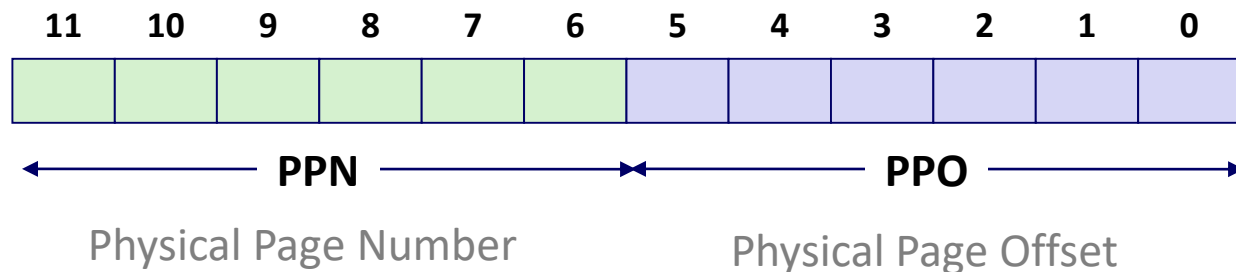
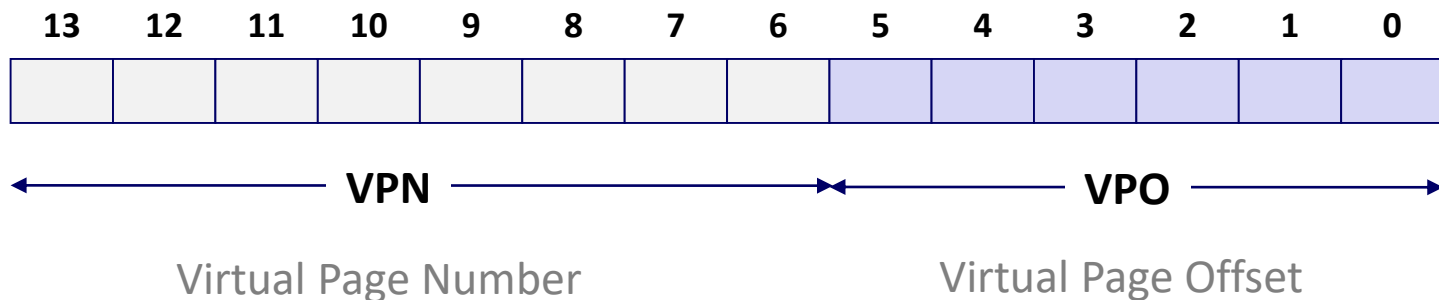
```
/* Sum rows is of n X n matrix a
   and store in vector b */
void sum_rows1(double *a, double *b, long n) {
    long i, j;
    for (j = 0; j < n; j++)
        for (i = 0; i < n; i++) {
            b[i] += a[i*n + j];
        }
}
```

```
/* Sum rows is of n X n matrix a
   and store in vector b */
void sum_rows2(double *a, double *b, long n) {
    long i, j;
    for (i = 0; i < n; i++) {
        double val = 0;
        for (j = 0; j < n; j++)
            val += a[i*n + j];
        b[i] = val;
    }
}
```

# Simple Memory System Example

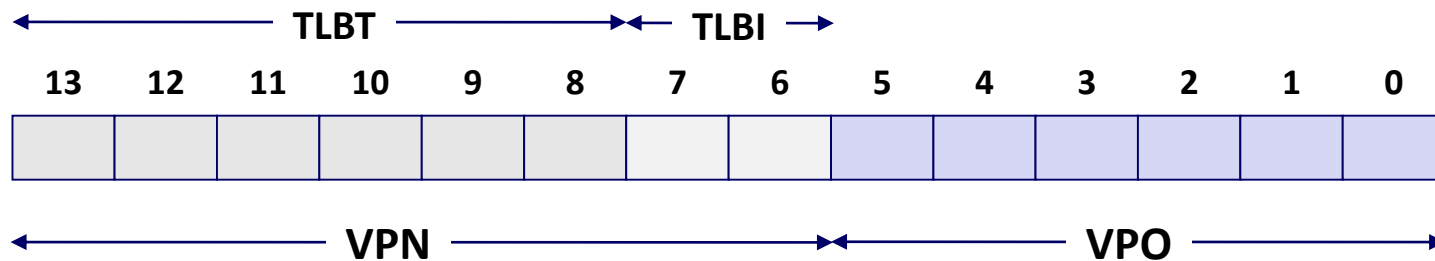
## ■ Addressing

- 14-bit virtual addresses
- 12-bit physical address
- Page size = 64 bytes



# 1. Simple Memory System TLB

- 16 entries
- 4-way associative



Set	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid
0	03	–	0	09	0D	1	00	–	0	07	02	1
1	03	2D	1	02	–	0	04	–	0	0A	–	0
2	02	–	0	08	–	0	06	–	0	03	–	0
3	07	–	0	03	0D	1	0A	34	1	02	–	0



## 2. Simple Memory System Page Table

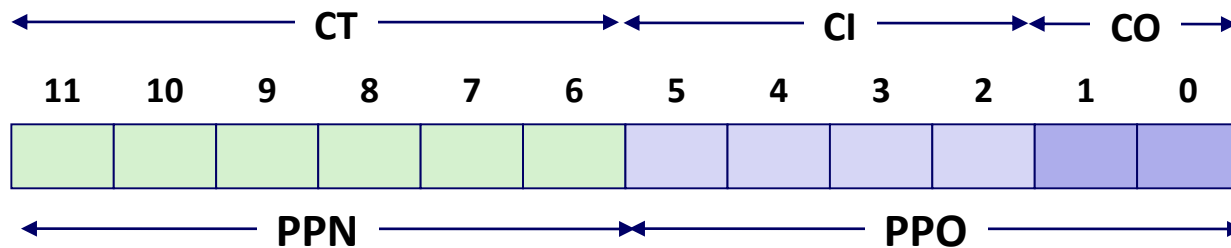
Only show first 16 entries (out of 256)

<i>VPN</i>	<i>PPN</i>	<i>Valid</i>
00	28	1
01	–	0
02	33	1
03	02	1
04	–	0
05	16	1
06	–	0
07	–	0

<i>VPN</i>	<i>PPN</i>	<i>Valid</i>
08	13	1
09	17	1
0A	09	1
0B	–	0
0C	–	0
0D	2D	1
0E	11	1
0F	0D	1

# 3. Simple Memory System Cache

- 16 lines, 4-byte block size
- Physically addressed
- Direct mapped

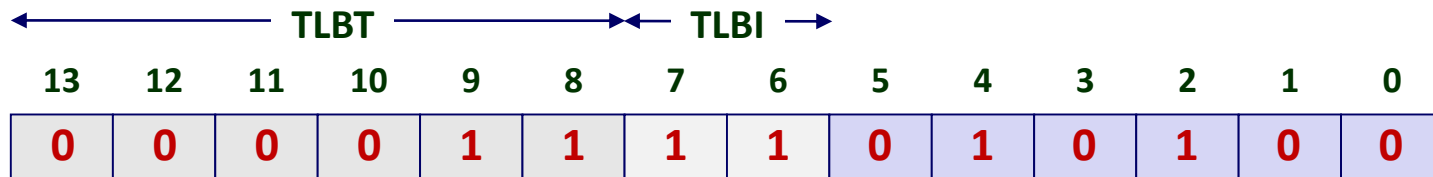


<i>Idx</i>	<i>Tag</i>	<i>Valid</i>	<i>B0</i>	<i>B1</i>	<i>B2</i>	<i>B3</i>
0	19	1	99	11	23	11
1	15	0	–	–	–	–
2	1B	1	00	02	04	08
3	36	0	–	–	–	–
4	32	1	43	6D	8F	09
5	0D	1	36	72	F0	1D
6	31	0	–	–	–	–
7	16	1	11	C2	DF	03

<i>Idx</i>	<i>Tag</i>	<i>Valid</i>	<i>B0</i>	<i>B1</i>	<i>B2</i>	<i>B3</i>
8	24	1	3A	00	51	89
9	2D	0	–	–	–	–
A	2D	1	93	15	DA	3B
B	0B	0	–	–	–	–
C	12	0	–	–	–	–
D	16	1	04	96	34	15
E	13	1	83	77	1B	D3
F	14	0	–	–	–	–

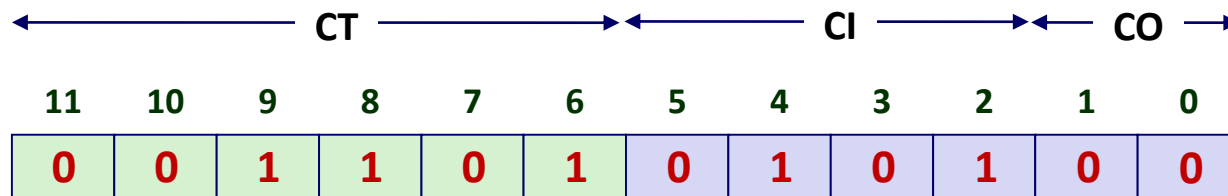
# Address Translation Example #1

Virtual Address: 0x03D4



VPN 0x0F    TLBI 0x3    TLBT 0x03    TLB Hit? Y    Page Fault? N    PPN: 0x0D

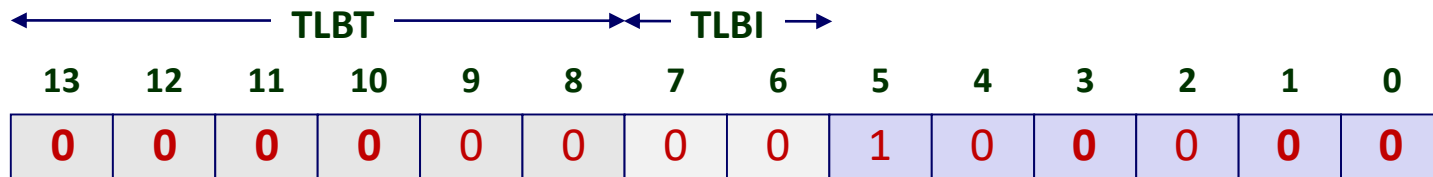
Physical Address



CO 0    CI 0x5    CT 0x0D    Hit? Y    Byte: 0x36

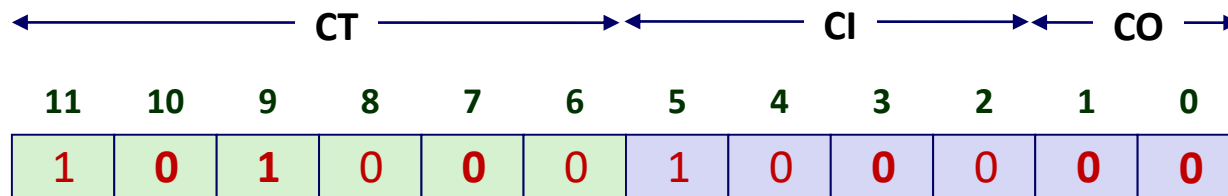
# Address Translation Example #2

Virtual Address: 0x0020



VPN 0x00 TLBI 0 TLBT 0x00 TLB Hit? N Page Fault? N PPN: 0x28

## Physical Address



CO 0 CI 0x8 CT 0x28 Hit? N Byte: Mem

# Hardest Test Questions?

## ■ Race Question

- Pthreads / Joining
- Forking / Shared Memory
- Semaphores

## ■ Attack Lab

- Stacks (you will draw a stack based on assembly)
- Construct / Identify an attack
- Test knowledge of datatypes

## ■ Maybe one or two more questions???

# Races and Semaphores

```
int a;
sem_t mutex[2];
int main() {
    Sem_init(&mutex[0], 0, 0);  /* mutex[0] = 0 */
    Sem_init(&mutex[1], 0, 1);  /* mutex[1] = 1 */
}
```

## Thread 1

```
sem_wait(&mutex[0]);
printf("what's up?");
sem_post(&mutex[1]);
```

## Thread 2

```
printf("hi");
sem_post(&mutex[0]);
sem_wait(&mutex[1]);
printf("not much");
```

## Thread 1

```
sem_wait(&mutex[0]);
printf("what's up?");
sem_post(&mutex[0]);
```

## Thread 2

```
printf("hi");
sem_post(&mutex[0]);
sem_wait(&mutex[0]);
printf("not much");
```

## Thread 1

```
sem_wait(&mutex[1]);
printf("what's up?");
sem_post(&mutex[1]);
```

## Thread 2

```
printf("hi");
sem_post(&mutex[1]);
sem_wait(&mutex[1]);
printf("not much");
```

# Stack and Data

```
int func(int i) {
    volatile union U {
        char s[16];
        int x;
    } u;

    u.x=0;
    u.s[i]=i;
    return u.x;
}
```

```
0x00000000004005d0 <+0>: movslq %edi,%rax
0x00000000004005d3 <+3>: movl  $0x0,-0x18(%rsp)
0x00000000004005db <+11>: mov  %dil,-0x18(%rsp,%rax,1)
0x00000000004005e0 <+16>: mov  -0x18(%rsp),%eax
0x00000000004005e4 <+20>: retq
```

```
func(0)=0x0
func(1)=0x100
func(2)=0x20000
func(3)=0x3000000
```

# x86-64 Linux Register Usage #1

## ■ **%rax**

- Return value
- Also caller-saved
- Can be modified by procedure

## ■ **%rdi, ..., %r9**

- Arguments
- Also caller-saved
- Can be modified by procedure

## ■ **%r10, %r11**

- Caller-saved
- Can be modified by procedure

Return value

**%rax**

Arguments

**%rdi**

**%rsi**

**%rdx**

**%rcx**

**%r8**

**%r9**

Caller-saved  
temporaries

**%r10**

**%r11**



# x86-64 Linux Register Usage

## #2

### ■ **%rbx, %r12, %r13, %r14**

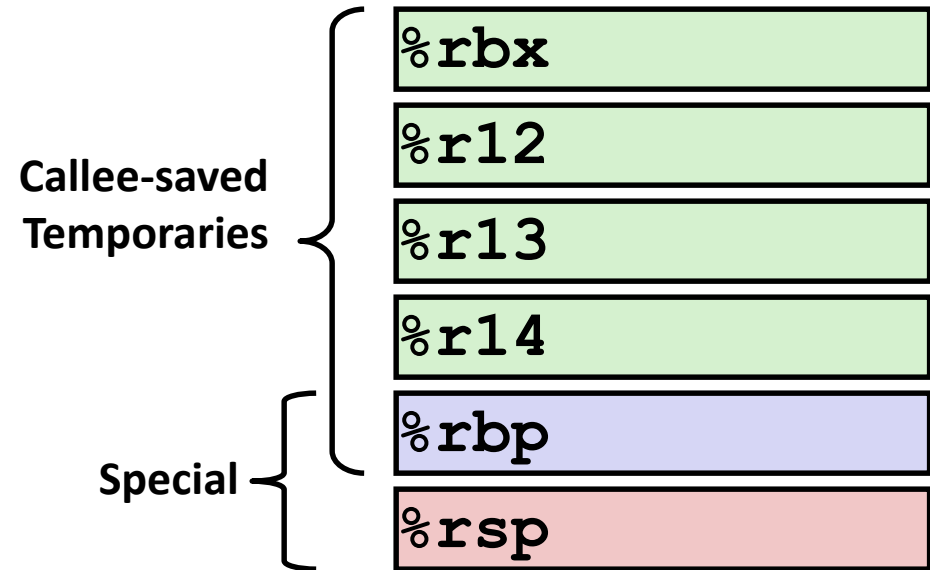
- Callee-saved
- Callee must save & restore

### ■ **%rbp**

- Callee-saved
- Callee must save & restore
- May be used as frame pointer
- Can mix & match

### ■ **%rsp**

- Special form of callee save
- Restored to original value upon exit from procedure



# Other Questions